




Movie and Music Recommendation System Using Machine Learning

MOVIE
Ricuripta
Alcouglcuin

DEMERY

Mioton

Project Details

Project Name

Movie and Music Recommendation System

Objective

To develop an intelligent recommendation system using machine learning that provides personalized movie and music suggestions to users, enhancing their content discovery experience.

Technologies Used

- Python
- Pandas
- Scikit-learn
- Cosine-Similarity
- Streamlit
- VS Code

Key Features

- Cosine Similarity Algorithm for content matching
- Content-Based Filtering for personalized recommendations
- Interactive User Interface (UI) built with Streamlit
- Real-time recommendation generation

Target Users

Individuals seeking personalized movie and music recommendations, as well as developers interested in learning about machine learning applications in recommendation systems.

Project Team

Student Team:

- DEBMALYA GHOSH (ROLL NO: 10301623026) - ELECTRICAL ENGINEERING
- PRIYANKA DUTTA (ROLL NO: 10300323104) - ELECTRONICS COMMUNICATION ENGINEERING
- PIJUSH MANNA (ROLL NO: 10301623053) - ELECTRICAL ENGINEERING

Group Name: TRIMIND

Faculty Advisor: MR. DEBASHIS SAHOO

Introduction: The Rise of Personalized Recommendations

The digital age has brought an unprecedented volume of content, making discovery a challenge for users. From endless streaming options to vast music libraries, the need for intelligent filtering has never been more critical. Personalized recommendation systems have emerged as essential tools, transforming how we interact with online platforms by suggesting content tailored to individual tastes.

Movie and music platforms, in particular, leverage these sophisticated systems to enhance user engagement, improve retention, and drive content consumption. By analysing past interactions and preferences, machine learning algorithms can predict what a user might enjoy next, turning a potentially overwhelming choice into a curated, enjoyable experience.



Enhanced User Experience

Streamlining content discovery to reduce decision fatigue.



Increased Engagement

Keeping users hooked with relevant and interesting suggestions.




Machine Learning Driven

Utilising powerful algorithms to predict individual preferences.

Understanding Recommendation Systems: Concepts and Types

Recommendation systems are broadly categorised into two primary types: content-based filtering and collaborative filtering. Content-based systems recommend items similar to those a user has previously expressed interest in, relying on features of the items themselves. For instance, if a user enjoys action movies, the system will suggest other action films with similar plot elements or directors.

Collaborative filtering, on the other hand, makes recommendations by leveraging the collective behaviour and preferences of a large group of users. It identifies users with similar tastes and recommends items that those 'similar' users have liked. Hybrid approaches combine the strengths of both methods, aiming to mitigate their individual weaknesses and provide more robust, accurate, and diverse recommendations. Central to many of these systems, particularly content-based ones, is cosine similarity, a metric used to quantify the likeness between two items based on their feature vectors.



The diagram above illustrates the fundamental differences between content-based and collaborative filtering, highlighting how hybrid approaches integrate both for superior results.

Dataset Selection and Feature Engineering

The quality of any recommendation system is heavily dependent on the richness and relevance of its underlying data. For music recommendations, datasets often include metadata such as genre, artist, track name, and various acoustic features like loudness, tempo, and speechiness. These attributes help differentiate tracks and establish similarity.

Similarly, movie datasets typically feature plot summaries, genres, cast information, director, and even user review sentiments. The challenge lies in transforming these disparate data types into a unified format that can be processed by machine learning algorithms. This involves combining textual features (which undergo cleaning, tokenization, and TF-IDF vectorization) with numerical features (which require normalisation) to create comprehensive feature vectors for each item. This meticulous preprocessing ensures that the system can accurately measure the similarity between different movies or songs.

Music Dataset Example

- Genre, Artist, Track Name
- Acoustic Features (e.g., Loudness, Speechiness)

Movie Dataset Example

- Plot Summaries, Genres, Cast
- User Reviews, Director

Preprocessing Steps

- Cleaning & Tokenization
- TF-IDF Vectorization for text
- Normalization for numerical features

Implementing Cosine Similarity for Recommendations

Cosine similarity is a fundamental metric in content-based recommendation systems, providing a measure of similarity between two non-zero vectors in an inner product space. Mathematically, it calculates the cosine of the angle between two vectors. A cosine value of 1 signifies identical direction (maximum similarity), 0 indicates orthogonality (no similarity), and -1 implies exactly opposite directions.

In our context, these vectors represent the feature sets of movies or songs. By computing the cosine similarity between a user's liked items and all other candidate items in the dataset, we can identify content that is most similar to their known preferences. The system then ranks these candidate items based on their similarity scores and presents the top-N recommendations. For instance, if a user enjoys a song with specific acoustic properties and genre tags, the system will recommend other songs whose feature vectors are closest in direction to that song's vector.

$$\cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|}$$

Where $\mathbf{A}$ and $\mathbf{B}$ are feature vectors, $\mathbf{A} \cdot \mathbf{B}$ is their dot product, and $\|\mathbf{A}\|$ and $\|\mathbf{B}\|$ are their magnitudes.

Building the Recommendation Engine in Python

Developing the core recommendation engine in Python involves leveraging powerful libraries tailored for data manipulation and machine learning. Pandas is indispensable for efficient data handling, enabling us to load, clean, and transform our music and movie datasets. Scikit-learn provides robust tools for text vectorization, specifically TF-IDF (Term Frequency-Inverse Document Frequency), and for calculating cosine similarity matrices.

The process begins with loading the datasets and performing the necessary preprocessing, as discussed earlier. Next, combined feature vectors are created for each item. The heart of the engine is the computation of the cosine similarity matrix, where each cell represents the similarity between two items. Finally, a function is defined to take an item (e.g., a movie title or song name) as input and return a list of top-N similar recommendations based on the pre-computed similarity matrix. Emphasising modular and reusable code ensures the system is scalable and maintainable.

Complete Python Implementation

This section provides the full Python code implementation, encompassing all steps from data loading and preprocessing to feature engineering, similarity computation, and defining the recommendation functions. This comprehensive script forms the core of both the movie and music recommendation systems discussed.

[Download Full Code \(Phttps://cdn.gamma.app/4lnbu3e76x2uf1q/31cc9b...](https://cdn.gamma.app/4lnbu3e76x2uf1q/31cc9b...)

Building the Recommendation Engine in Python

With the data prepared and cosine similarity chosen as our metric, the next crucial step is to construct the recommendation engine itself. This involves transforming the textual features into numerical representations, refining these representations, calculating similarity scores, and finally, developing the functions that deliver personalised recommendations. Python's rich ecosystem of libraries, particularly `scikit-learn` and `nlTK`, proves invaluable in orchestrating this process efficiently.

Our approach ensures that both movie and music recommendation systems leverage a consistent methodology, making the engine robust and adaptable. By breaking down the task into logical, sequential steps, we create a clear pathway from raw text data to actionable recommendations.

1

Feature Extraction & Vectorization

The first step involves converting the combined 'tags' (overview, genres, keywords, cast, crew for movies; name, artist, genre for music) into a numerical format using `CountVectorizer`. This technique counts the occurrences of each word, creating a vector for every item. We limit features to 5,000 and remove common English stop words to focus on meaningful terms.

2

Text Preprocessing with Stemming

To enhance the accuracy of similarity calculations, we apply stemming using `PorterStemmer` from `nlTK`. Stemming reduces words to their root form (e.g., 'loving', 'loved', 'loves' all become 'love'), ensuring that variations of the same word are treated identically, thus improving the identification of thematic similarities.

3

Cosine Similarity Calculation

Once the text features are vectorised and stemmed, `cosine_similarity` from `scikit-learn` is employed. This calculates a score between 0 and 1 for every pair of items, indicating how similar they are based on their tag vectors. A higher score signifies greater similarity.

4

Recommendation Logic

Implementation

Dedicated Python functions, `recommend_movie` and `recommend_music`, are built to utilise these similarity scores. Given an input item, they retrieve its index, access its similarity scores, sort them in descending order, and return the top 5 most similar items, excluding the input itself.

5

Model Persistence

Finally, the processed dataframes and the computed similarity matrices are saved using the `pickle` module. This allows the trained models and data artefacts to be quickly loaded for deployment, avoiding the need to re-process everything each time the application runs.

Deploying with Streamlit in VS Code: Interactive User Interface

To make our recommendation engine accessible and interactive, Streamlit provides an excellent, lightweight framework for rapid web application development. Its simplicity allows developers to turn Python scripts into interactive web apps with minimal effort, making it ideal for prototyping and demonstrating machine learning models.

Integrating our recommendation logic with Streamlit involves creating a user-friendly interface in VS Code. This interface will incorporate various Streamlit widgets, such as text input boxes for searching movie or song titles, and dropdown menus for selecting genres or artists. When a user provides input, the Streamlit app will call our Python recommendation function, process the request, and dynamically display the top recommendations, complete with titles, and potentially images or other relevant metadata. The seamless integration within the VS Code environment allows for efficient development, testing, and debugging, bringing the recommendation system to life.

Streamlit Features

- Rapid web app development
- Interactive widgets (sliders, text input)
- Dynamic content display
- Seamless VS Code integration

Key Integration Steps

- Import Streamlit
- Define user input widgets
- Call recommendation function
- Display results using `st.write()`

Evaluation and Performance Metrics

Evaluating the performance of a recommendation system is crucial for understanding its effectiveness and identifying areas for improvement. Accuracy and relevance are often measured through user feedback, but more quantitative assessments utilise offline metrics such as precision, recall, and F1-score. These metrics help gauge how many of the recommended items are truly relevant to the user's preferences.

However, recommendation systems face several inherent challenges. The 'cold start problem' occurs when there is insufficient data for new users or new items, making it difficult to generate accurate recommendations. Data sparsity, where user-item interaction matrices are mostly empty, can also hinder performance. Overfitting to popular items can lead to a lack of diversity in recommendations. Strategies to combat these issues include employing hybrid models, regularly retraining the system with fresh data, and implementing sophisticated user profiling techniques. A music recommendation system, for instance, once achieved approximately 90% accuracy by combining content-based and collaborative filtering methods, demonstrating the power of hybrid approaches.

90%

Approximate accuracy achieved by a hybrid music recommendation system.

3

Common challenges: cold start, data sparsity, popularity bias.

4

Strategies for improvement: hybrid models, retraining, profiling.

Real-World Applications and Future Enhancements

The impact of recommendation systems is ubiquitous in our daily digital lives. Giants like Spotify leverage them for features such as "Discover Weekly," which curates personalized playlists, and Netflix employs them to offer tailored movie and TV show suggestions, significantly influencing user content consumption. These systems have moved beyond simple suggestions to become integral parts of the user experience, often dictating what we see and hear next.

Future enhancements in recommendation technology are poised to incorporate more advanced machine learning techniques. Deep learning and Natural Language Processing (NLP) can enable richer content understanding, for example, by performing sentiment analysis on movie reviews or extracting nuanced thematic elements from song lyrics. Adding explainability to recommendations, perhaps through SHAP values, can increase user trust by showing why a particular item was suggested. Furthermore, the expansion to cross-domain recommendations, combining music and movies for mood-based experiences, presents exciting possibilities for even more immersive and personalised discovery journeys.

Conclusion: Empowering Discovery Through Machine Learning

This document has outlined the journey of building a practical movie and music recommendation system, leveraging the power of machine learning, specifically cosine similarity, and deploying it with an interactive user interface using Streamlit in VSCode. We began by understanding the critical need for personalized recommendations in today's content-rich environment, explored the fundamental concepts of recommendation types, and delved into the intricacies of dataset selection and feature engineering.

The core of our system relies on the mathematical elegance of cosine similarity to quantify item likeness, translated into a functional Python engine. Finally, we demonstrated how Streamlit can transform this engine into an engaging, interactive application. The journey from raw data to personalised suggestions highlights the importance of combining technical rigour with user-centric design. We encourage further experimentation and extension of this system with advanced machine learning techniques, as personalized recommendations continue to evolve, transforming passive consumption into truly engaging discovery journeys.